

Homework 4

July 15, 2025

Author:Alejandro M.Ouslan

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.formula.api as smf
import statsmodels.stats.api as sms
from statsmodels.tsa.arima.model import ARIMA
import statsmodels.api as sm
from statsmodels.tsa.ar_model import AutoReg
import numpy as np
from statsmodels.stats.diagnostic import acorr_ljungbox
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
[2]: df = pd.read_excel("data/MONEYDEM-1.xls")
df['year'] = df['DATE'].astype(int)
df['quarter'] = ((df['DATE'] - df['year']) * 10).round().astype(int)
df['date'] = df['year'].astype(str) + 'Q' + df['quarter'].astype(str)
df['date'] = pd.PeriodIndex(df['date'], freq='Q')
df.set_index('date', inplace=True)
df
```

```
[2]:
```


	DATE	TB3mo	TB1yr	year	quarter
date					
1959Q3	1959.3	3.540000	4.493333	1959	3
1959Q4	1959.4	4.230000	4.740000	1959	4
1960Q1	1960.1	3.873333	4.360000	1960	1
1960Q2	1960.2	2.993333	3.646667	1960	2
1960Q3	1960.3	2.360000	2.903333	1960	3
...
2000Q1	2000.1	5.520000	5.816667	2000	1
2000Q2	2000.2	5.713333	5.856667	2000	2
2000Q3	2000.3	6.016667	5.803333	2000	3
2000Q4	2000.4	6.016667	5.630000	2000	4
2001Q1	2001.1	4.816667	4.416667	2001	1

[167 rows x 5 columns]


1 Problems

1.1 Problem a

```
[3]: #1a
plt.figure(figsize=(8, 5))
plt.plot(df.index.to_timestamp(), df['TB3mo'], marker='o')
plt.title('Quarterly Data')
plt.xlabel('Date')
plt.ylabel('TB3mo')
plt.grid(True)
plt.tight_layout()
plt.show()
```




```
[4]: plt.figure(figsize=(8, 5))
plt.plot(df.index.to_timestamp(), df['TB1yr'], marker='o')
plt.title('Quarterly Data')
plt.xlabel('Date')
plt.ylabel('TB1yr')
plt.grid(True)
plt.tight_layout()
plt.show()
```



- No, they appear to have changing mean and variance given that in 1980 had a vilont spike

1.2 Problem b

```
[5]: plt.figure(figsize=(8, 5))
plt.plot(df.index.to_timestamp(), df['TB1yr'], marker='o')
plt.plot(df.index.to_timestamp(), df['TB3mo'], marker='o')
plt.title('Quarterly Data')
plt.xlabel('Date')
plt.grid(True)
plt.tight_layout()
plt.show()
```



- They appear to be practically identical the seam to follow the same trend

1.3 Problem c

```
[6]: model = smf.ols("TB1yr ~ TB3mo", data=df).fit()
print(model.summary())
```

OLS Regression Results

=====						
Dep. Variable:	TB1yr		R-squared:	0.980		
Model:	OLS		Adj. R-squared:	0.980		
Method:	Least Squares		F-statistic:	7975.		
Date:	Tue, 15 Jul 2025		Prob (F-statistic):	1.30e-141		
Time:	21:30:14		Log-Likelihood:	-56.689		
No. Observations:	167		AIC:	117.4		
Df Residuals:	165		BIC:	123.6		
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	0.6981	0.067	10.486	0.000	0.567	0.830
TB3mo	0.9167	0.010	89.302	0.000	0.896	0.937
=====						
Omnibus:	15.536		Durbin-Watson:	0.623		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	39.846		

Skew:	0.289	Prob(JB):	2.23e-09
Kurtosis:	5.322	Cond. No.	16.7

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1.4 Problem d

- A 1 percentage point increase in the 3-month Treasury rate (short-term) is associated with a 0.9167 percentage point increase in the 1-year Treasury rate (long-term), on average, holding other factors constant.

1.5 Problem e

```
[7]: t_test_result = model.t_test('TB3mo = 1')
      print(t_test_result)
```

Test for Constraints						
	coef	std err	t	P> t	[0.025	0.975]
c0	0.9167	0.010	-8.112	0.000	0.896	0.937


=====

- since p-value is 0.000 there is strong but less-than-perfect pass-through from short- to long-term rates.

1.6 Problem f

```
[8]: residuals = model.resid

plt.figure(figsize=(8, 5))
plt.scatter(df['TB3mo'], residuals, alpha=0.7)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('TB3mo (3-month Treasury rate)')
plt.ylabel('Residuals')
plt.title('Residuals vs TB3mo')
plt.grid(True)
plt.show()
```



- No there does not seem to be a pattern

1.7 Problem g

```
[9]: residuals = model.resid
     exog = model.model.exog

     white_test = sms.het_white(residuals, exog)

     lm_stat, lm_pvalue, f_stat, f_pvalue = white_test

     print(f"White test LM statistic: {lm_stat:.4f}")
     print(f"White test LM p-value: {lm_pvalue:.4f}")
     print(f"White test F statistic: {f_stat:.4f}")
     print(f"White test F p-value: {f_pvalue:.4f}")
```

```
White test LM statistic: 31.3738
White test LM p-value: 0.0000
White test F statistic: 18.9687
White test F p-value: 0.0000
```

- given that the p-value is 0.0000 there is evidence of heteroskedasticity

1.8 Problem h

```
[10]: model_robust = smf.ols("TB1yr ~ TB3mo", data=df).fit(cov_type='HCO')
print(model_robust.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          TB1yr      R-squared:                0.980
Model:                  OLS        Adj. R-squared:           0.980
Method:                 Least Squares    F-statistic:         3085.
Date:                  Tue, 15 Jul 2025    Prob (F-statistic):    1.05e-108
Time:                  21:30:14      Log-Likelihood:        -56.689
No. Observations:      167          AIC:                   117.4
Df Residuals:          165          BIC:                   123.6
Df Model:               1
Covariance Type:       HCO
=====
                        coef      std err          z      P>|z|      [0.025      0.975]
-----
Intercept              0.6981      0.088       7.940      0.000      0.526      0.870
TB3mo                  0.9167      0.017     55.539      0.000      0.884      0.949
=====
Omnibus:               15.536    Durbin-Watson:           0.623
Prob(Omnibus):          0.000    Jarque-Bera (JB):        39.846
Skew:                   0.289    Prob(JB):                 2.23e-09
Kurtosis:               5.322    Cond. No.                 16.7
=====
```

Notes:

[1] Standard Errors are heteroscedasticity robust (HCO)

- The coefecients did not change, however the std err did increase for both the intercept and the coefficient When heteroskedasticity is present, regular SEs underestimate the true variability of the coefficients, so the robust SEs tend to be larger and more reliable.

1.9 Problem j

```
[11]: df['D'] = (df['TB3mo'] > 10.00).astype(int)
model_with_dummy = smf.ols("TB1yr ~ TB3mo + D", data=df).fit()
print(model_with_dummy.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          TB1yr      R-squared:                0.981
Model:                  OLS        Adj. R-squared:           0.981
Method:                 Least Squares    F-statistic:         4176.
Date:                  Tue, 15 Jul 2025    Prob (F-statistic):    2.16e-141
Time:                  21:30:14      Log-Likelihood:        -52.401
```

No. Observations: 167 AIC: 110.8
Df Residuals: 164 BIC: 120.2
Df Model: 2
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.5551	0.081	6.832	0.000	0.395	0.716
TB3mo	0.9452	0.014	67.730	0.000	0.918	0.973
D	-0.4456	0.152	-2.940	0.004	-0.745	-0.146
Omnibus:	15.483		Durbin-Watson:	0.597		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	26.029		
Skew:	0.476		Prob(JB):	2.23e-06		
Kurtosis:	4.684		Cond. No.	40.8		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1.10 Problem k

- given that the p-value of delta is 0.004 the dummy is revelat.

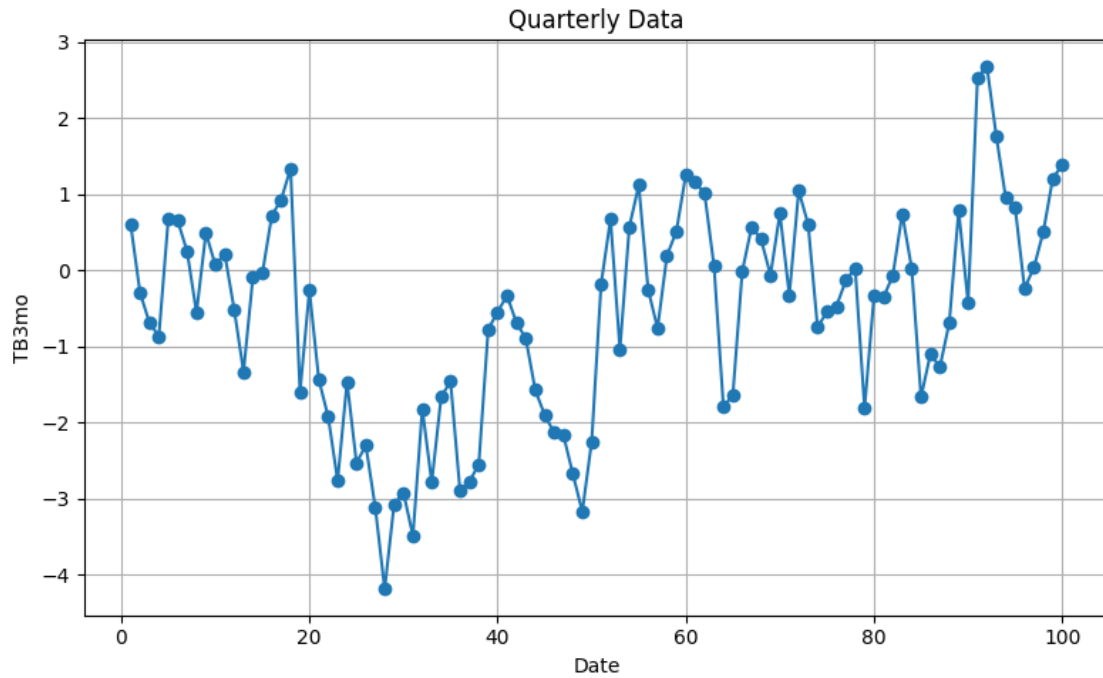
1.11 Problem l

- The coefien increased but also did the standard error

2 Problem 2

2.1 Problem A


```
[12]: df = pd.read_excel("data/SIM_2-1.xls")
plt.figure(figsize=(8, 5))
plt.plot(df["OBS"], df['Y1'], marker='o')
plt.title('Quarterly Data')
plt.xlabel('Date')
plt.ylabel('TB3mo')
plt.grid(True)
plt.tight_layout()
plt.show()
```





- given that the series has not broke out dwon or up and it seems to hover aroun 1 we could say that it is stationary

2.2 Problem B

```
[13]: fig = plot_acf(df['Y1'], lags=10)
plt.show()
plot_pacf(df['Y1'], lags=15)
plt.show()
```





- given that the series has not broke out dwon or up and it seems to hover around 1 we could say that it is stationary

2.3 Problem C

```
[14]: # AR(1)
res = AutoReg(df['Y1'], lags =1).fit()
print(res.summary())
y_true = res.model.endog[res.model._hold_back:]
y_pred = res.fittedvalues

ssr = np.sum((y_true - y_pred) ** 2)

tss = np.sum((y_true - np.mean(y_true)) ** 2)

n = len(y_true)
k = res.df_model + 1

r2 = 1 - ssr / tss
r2_adj = 1 - (ssr / (n - k)) / (tss / (n - 1))
```

```
print("R^2:", round(r2, 4))
print("Adjusted R^2:", round(r2_adj, 4))
```

AutoReg Model Results

```
=====
Dep. Variable:          Y1      No. Observations:          100
Model:                AutoReg(1)  Log Likelihood          -132.061
Method:              Conditional MLE  S.D. of innovations          0.919
Date:                Tue, 15 Jul 2025  AIC                270.122
Time:                21:30:14  BIC                277.907
Sample:              1      HQIC                273.272
                    100
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-0.1372	0.100	-1.366	0.172	-0.334	0.060
Y1.L1	0.7545	0.067	11.269	0.000	0.623	0.886

Roots

	Real	Imaginary	Modulus	Frequency
AR.1	1.3254	+0.0000j	1.3254	0.0000

R^2: 0.5619

Adjusted R^2: 0.5528

```
[15]: # AR(2)
res = AutoReg(df['Y1'], lags=2).fit()
print(res.summary())
y_true = res.model.endog[res.model._hold_back:]
y_pred = res.fittedvalues

ssr = np.sum((y_true - y_pred) ** 2)

tss = np.sum((y_true - np.mean(y_true)) ** 2)

n = len(y_true)
k = res.df_model + 1

r2 = 1 - ssr / tss
r2_adj = 1 - (ssr / (n - k)) / (tss / (n - 1))

print("R^2:", round(r2, 4))
print("Adjusted R^2:", round(r2_adj, 4))
```

AutoReg Model Results

Dep. Variable:	Y1	No. Observations:	100
Model:	AutoReg(2)	Log Likelihood	-130.629
Method:	Conditional MLE	S.D. of innovations	0.918
Date:	Tue, 15 Jul 2025	AIC	269.258
Time:	21:30:14	BIC	279.598
Sample:	2	HQIC	273.440
	100		

	coef	std err	z	P> z	[0.025	0.975]
const	-0.1144	0.102	-1.116	0.264	-0.315	0.086
Y1.L1	0.6939	0.101	6.896	0.000	0.497	0.891
Y1.L2	0.0870	0.101	0.861	0.389	-0.111	0.285

Roots

	Real	Imaginary	Modulus	Frequency
AR.1	1.2462	+0.0000j	1.2462	0.0000
AR.2	-9.2181	+0.0000j	9.2181	0.5000

R²: 0.567

Adjusted R²: 0.5532

```
[16]: # ARMA(1,1)
arma_mod = ARIMA(df['Y1'], order=(1, 1, 0)).fit()
print(arma_mod.summary())

y_true = df['Y1'].diff().values

y_pred = arma_mod.fittedvalues

ssr = np.sum((y_true - y_pred) ** 2)
tss = np.sum((y_true - np.mean(y_true)) ** 2)

n = len(y_true)
k = arma_mod.df_model

r2 = 1 - ssr / tss
r2_adj = 1 - (ssr / (n - k)) / (tss / (n - 1))

print("R^2:", round(r2, 4))
print("Adjusted R^2:", round(r2_adj, 4))

print("AIC:", arma_mod.aic)
print("BIC:", arma_mod.bic)
```

SARIMAX Results

```

=====
Dep. Variable:          Y1      No. Observations:          100
Model:                 ARIMA(1, 1, 0)      Log Likelihood          -136.382
Date:                 Tue, 15 Jul 2025      AIC              276.764
Time:                 21:30:14      BIC              281.954
Sample:                 0      HQIC              278.864
                        - 100
Covariance Type:          opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.1982	0.106	-1.868	0.062	-0.406	0.010
sigma2	0.9203	0.125	7.356	0.000	0.675	1.165

```

=====
===
Ljung-Box (L1) (Q):          0.05      Jarque-Bera (JB):
0.23
Prob(Q):                    0.82      Prob(JB):
0.89
Heteroskedasticity (H):      1.04      Skew:
-0.01
Prob(H) (two-sided):        0.91      Kurtosis:
3.23
=====
===

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
R^2: nan
Adjusted R^2: nan
AIC: 276.7637125558847
BIC: 281.95395225615385

```

```

[17]: # ARMA(1,4)
arma_mod = ARIMA(df['Y1'], order=(1, 4, 0)).fit()
print(arma_mod.summary())
y_true = df['Y1'].diff().values

y_pred = arma_mod.fittedvalues
ssr = np.sum((y_true - y_pred) ** 2)
tss = np.sum((y_true - np.mean(y_true)) ** 2)

n = len(y_true)
k = arma_mod.df_model

r2 = 1 - ssr / tss

```

```

r2_adj = 1 - (ssr / (n - k)) / (tss / (n - 1))

print("R^2:", round(r2, 4))
print("Adjusted R^2:", round(r2_adj, 4))

print("AIC:", arma_mod.aic)
print("BIC:", arma_mod.bic)

```

SARIMAX Results

```

=====
Dep. Variable:          Y1      No. Observations:          100
Model:                ARIMA(1, 4, 0)      Log Likelihood      -248.774
Date:                 Tue, 15 Jul 2025      AIC                501.548
Time:                 21:30:14      BIC                506.677
Sample:               0      HQIC                503.621
                   - 100
Covariance Type:      opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.7572	0.064	-11.873	0.000	-0.882	-0.632
sigma2	10.3403	1.580	6.546	0.000	7.244	13.436

```

=====
Ljung-Box (L1) (Q):          23.50      Jarque-Bera (JB):
2.57
Prob(Q):                     0.00      Prob(JB):
0.28
Heteroskedasticity (H):      1.16      Skew:
0.39
Prob(H) (two-sided):        0.68      Kurtosis:
2.86
=====

```

```

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
R^2: nan
Adjusted R^2: nan
AIC: 501.54819826952695
BIC: 506.67689465246264

```

```

[18]: # ARMA(2,1)
arma_mod = ARIMA(df['Y1'], order=(2, 1, 0)).fit()
print(arma_mod.summary())
y_true = df['Y1'].diff().values

```

```

y_pred = arma_mod.fittedvalues
ssr = np.sum((y_true - y_pred) ** 2)
tss = np.sum((y_true - np.mean(y_true)) ** 2)

n = len(y_true)
k = arma_mod.df_model

r2 = 1 - ssr / tss
r2_adj = 1 - (ssr / (n - k)) / (tss / (n - 1))

print("R^2:", round(r2, 4))
print("Adjusted R^2:", round(r2_adj, 4))

print("AIC:", arma_mod.aic)
print("BIC:", arma_mod.bic)

```

SARIMAX Results

```

=====
Dep. Variable:          Y1      No. Observations:          100
Model:                ARIMA(2, 1, 0)  Log Likelihood        -135.770
Date:                Tue, 15 Jul 2025  AIC                    277.540
Time:                21:30:14    BIC                    285.325
Sample:              0      HQIC                    280.690
                        - 100
Covariance Type:      opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.2212	0.109	-2.034	0.042	-0.434	-0.008
ar.L2	-0.1107	0.105	-1.052	0.293	-0.317	0.095
sigma2	0.9087	0.122	7.439	0.000	0.669	1.148

```

=====
Ljung-Box (L1) (Q):          0.01  Jarque-Bera (JB):
0.45
Prob(Q):                    0.92  Prob(JB):
0.80
Heteroskedasticity (H):      1.01  Skew:
0.05
Prob(H) (two-sided):        0.97  Kurtosis:
3.31
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

R²: nan
Adjusted R²: nan
AIC: 277.5398840453032
BIC: 285.325243595707

2.4 Problem D

```
[19]: #AR(2)
res = AutoReg(df['Y1'], lags=2, trend="n").fit()
print(res.summary())
y_true = df['Y1'].diff().values

y_pred = arma_mod.fittedvalues
ssr = np.sum((y_true - y_pred) ** 2)
tss = np.sum((y_true - np.mean(y_true)) ** 2)

n = len(y_true)
k = arma_mod.df_model
r2 = 1 - ssr / tss
r2_adj = 1 - (ssr / (n - k)) / (tss / (n - 1))

print("R^2:", round(r2, 4))
print("Adjusted R^2:", round(r2_adj, 4))

print("AIC:", arma_mod.aic)
print("BIC:", arma_mod.bic)
```

AutoReg Model Results

```
=====
Dep. Variable:          Y1    No. Observations:          100
Model:                  AutoReg(2)    Log Likelihood          -131.248
Method:                  Conditional MLE    S.D. of innovations          0.923
Date:                    Tue, 15 Jul 2025    AIC          268.496
Time:                    21:30:14    BIC          276.251
Sample:                  2    HQIC          271.633
                        100
=====
```

	coef	std err	z	P> z	[0.025	0.975]
Y1.L1	0.7102	0.100	7.087	0.000	0.514	0.907
Y1.L2	0.1051	0.100	1.046	0.295	-0.092	0.302

Roots

	Real	Imaginary	Modulus	Frequency
AR.1	1.1963	+0.0000j	1.1963	0.0000
AR.2	-7.9532	+0.0000j	7.9532	0.5000

R²: nan
Adjusted R²: nan
AIC: 277.5398840453032
BIC: 285.325243595707

```
[20]: # ARIMA(1,1)
arma_mod = ARIMA(df['Y1'], order=(1, 1, 0), trend="n").fit()
print(arma_mod.summary())
y_true = df['Y1'].diff().values

y_pred = arma_mod.fittedvalues

ssr = np.sum((y_true - y_pred) ** 2)
tss = np.sum((y_true - np.mean(y_true)) ** 2)

n = len(y_true)
k = arma_mod.df_model

r2 = 1 - ssr / tss
r2_adj = 1 - (ssr / (n - k)) / (tss / (n - 1))

print("R2:", round(r2, 4))
print("Adjusted R2:", round(r2_adj, 4))

print("AIC:", arma_mod.aic)
print("BIC:", arma_mod.bic)
```

SARIMAX Results

```
=====
Dep. Variable:          Y1      No. Observations:          100
Model:                ARIMA(1, 1, 0)      Log Likelihood          -136.382
Date:                 Tue, 15 Jul 2025      AIC              276.764
Time:                 21:30:14      BIC              281.954
Sample:               0      HQIC              278.864
                  - 100
Covariance Type:      opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.1982	0.106	-1.868	0.062	-0.406	0.010
sigma2	0.9203	0.125	7.356	0.000	0.675	1.165

```
=====
Ljung-Box (L1) (Q):          0.05      Jarque-Bera (JB):
0.23
Prob(Q):                     0.82      Prob(JB):
0.89
Heteroskedasticity (H):      1.04      Skew:
```

```

-0.01
Prob(H) (two-sided):          0.91    Kurtosis:
3.23
=====
===

```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
```

```
R^2: nan
```

```
Adjusted R^2: nan
```

```
AIC: 276.7637125558847
```

```
BIC: 281.95395225615385
```

2.5 Problem E

- for the part c looking at the AIC the best model is the AR(2) given that it has the smallest AIC fro the part ed the best is sitll AR(2) given that it has the smallest AIC

2.6 Problem F

- yes because looking at the AIC it shouldgest that the best model is AR(2) given that the simulated model was created with AR(1) one would expect AR(1) would be the best model

2.7 Problem G

```

[21]: # AR(2)
model_ar2 = AutoReg(df['Y1'], lags=2, old_names=False)
res_ar2 = model_ar2.fit()

residuals = res_ar2.resid

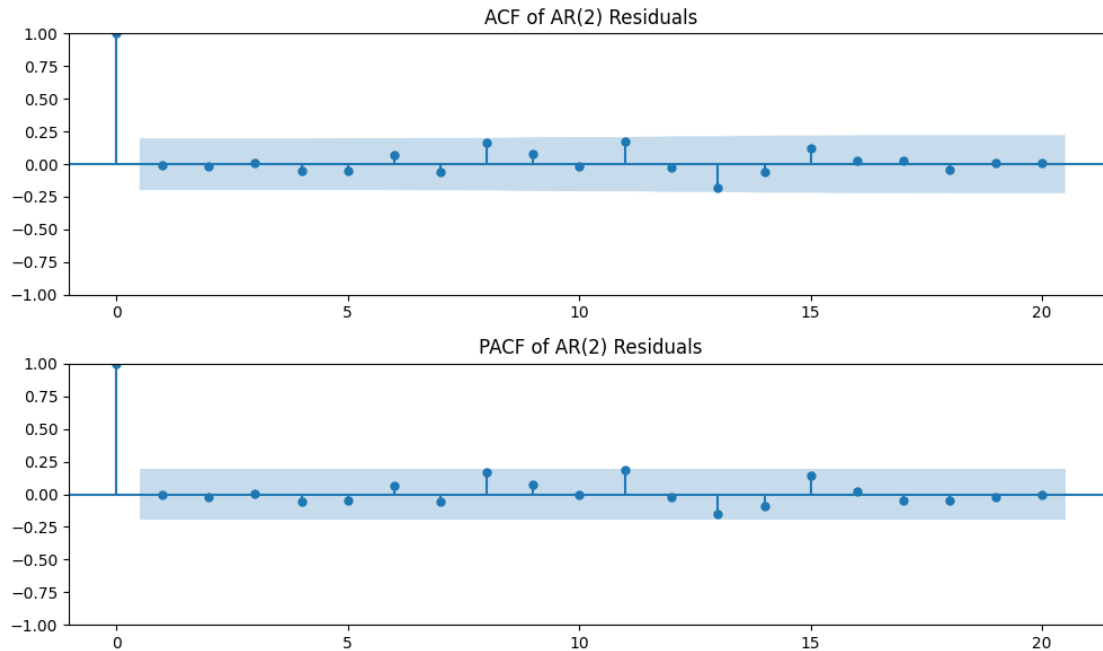
fig, ax = plt.subplots(2, 1, figsize=(10, 6))

plot_acf(residuals, ax=ax[0], lags=20)
ax[0].set_title('ACF of AR(2) Residuals')

plot_pacf(residuals, ax=ax[1], lags=20, method='ywm')
ax[1].set_title('PACF of AR(2) Residuals')

plt.tight_layout()
plt.show()

```



- yes they look like random noise

3 Problem

3.1 Problem A


```
[22]: df = pd.read_excel("data/QUARTERLY-1.xls")
df['date'] = pd.PeriodIndex(df['Date'], freq='Q')
df.set_index('date', inplace=True)
df = df[["CPINSA", "Date"]]
df
```

```
[22]:
```

	CPINSA	Date
date		
1960Q1	30.57	1960Q1
1960Q2	30.60	1960Q2
1960Q3	30.60	1960Q3
1960Q4	30.80	1960Q4
1961Q1	30.80	1961Q1
...
2007Q1	209.01	2007Q1
2007Q2	210.37	2007Q2
2007Q3	211.16	2007Q3
2007Q4	212.37	2007Q4
2008Q1	213.96	2008Q1

[193 rows x 2 columns]


```
[23]: plt.figure(figsize=(8, 5))
plt.plot(df.index.to_timestamp(), df['CPINSA'], marker='o')
plt.title('Quarterly Data')
plt.xlabel('Date')
plt.ylabel('TB3mo')
plt.grid(True)
plt.tight_layout()
plt.show()
```




- does not look stationary

3.2 Problem B


```
[24]: fig = plot_acf(df['CPINSA'], lags=25)
plt.show()
plot_pacf(df['CPINSA'], lags=25)
plt.show()
```





3.3 Problem C


```
[25]: df["log_CPINSA"] = np.log((df["CPINSA"] / df["CPINSA"].shift(1)))
plt.figure(figsize=(8, 5))
plt.plot(df.index.to_timestamp(), df['log_CPINSA'], marker='o')
plt.title('Quarterly Data')
plt.xlabel('Date')
plt.ylabel('TB3mo')
plt.grid(True)
plt.tight_layout()
plt.show()
```




- There was a dip but it seems to be stationary given that there is no strong trend

3.4 Problem D


```
[26]: fig = plot_acf(df['log_CPINSA'].dropna(), lags=25)
plt.show()
plot_pacf(df['log_CPINSA'].dropna(), lags=25)
plt.show()
```



3.5 Problem E


```
[27]: df["log_CPINSA4"] = np.log((df["CPINSA"] / df["CPINSA"].shift(4)))
plt.figure(figsize=(8, 5))
plt.plot(df.index.to_timestamp(), df['log_CPINSA4'], marker='o')
plt.title('Quarterly Data')
plt.xlabel('Date')
plt.ylabel('TB3mo')
plt.grid(True)
plt.tight_layout()
plt.show()
```




- there was a dip but it seems to be stationary given that there is no strong trend

3.6 Problem F

```
[28]: fig = plot_acf(df['log_CPINSA4'].dropna(), lags=25)
plt.show()
fig = plot_pacf(df['log_CPINSA4'].dropna(), lags=25)
plt.show()
```





3.7 Problem G

```
[29]: res = AutoReg(df['log_CPINSA4'].dropna(), lags =5).fit()
print(res.summary())
```

```

AutoReg Model Results
=====
Dep. Variable:          log_CPINSA4    No. Observations:          189
Model:                 AutoReg(5)     Log Likelihood              743.312
Method:                Conditional MLE  S.D. of innovations         0.004
Date:                  Tue, 15 Jul 2025  AIC                          -1472.624
Time:                  21:30:15         BIC                          -1450.120
Sample:                06-30-1962      HQIC                         -1463.503
                   - 03-31-2008
=====
==

```

	coef	std err	z	P> z	[0.025
0.975]					
const	0.0012	0.001	1.826	0.068	-8.5e-05

```

-----
--

```

```

0.002
log_CPINSA4.L1      1.4432      0.068      21.189      0.000      1.310
1.577
log_CPINSA4.L2      -0.4339      0.117      -3.700      0.000      -0.664
-0.204
log_CPINSA4.L3       0.2177      0.120       1.811      0.070      -0.018
0.453
log_CPINSA4.L4      -0.6346      0.117      -5.425      0.000      -0.864
-0.405
log_CPINSA4.L5       0.3809      0.068       5.621      0.000       0.248
0.514

```

Roots				
	Real	Imaginary	Modulus	Frequency
AR.1	-0.7595	-1.0363j	1.2848	-0.3507
AR.2	-0.7595	+1.0363j	1.2848	0.3507
AR.3	1.0440	-0.0000j	1.0440	-0.0000
AR.4	1.0704	-0.6143j	1.2342	-0.0829
AR.5	1.0704	+0.6143j	1.2342	0.0829

```

[30]: y_true = res.model.endog[res.model._hold_back:]
      y_pred = res.fittedvalues

      ssr = np.sum((y_true - y_pred) ** 2)

      tss = np.sum((y_true - np.mean(y_true)) ** 2)

      n = len(y_true)
      k = res.df_model + 1

      r2 = 1 - ssr / tss
      r2_adj = 1 - (ssr / (n - k)) / (tss / (n - 1))

      print("R^2:", round(r2, 4))
      print("Adjusted R^2:", round(r2_adj, 4))

```

```

R^2: 0.9708
Adjusted R^2: 0.9698

```

```

[31]: arma_mod = ARIMA(df['log_CPINSA4'].dropna(), order=(0, 0, 10)).fit()
      print(arma_mod.summary())

```

```

/home/ouslan/Documents/Github/ECON-124/.venv/lib/python3.10/site-
packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible
starting MA parameters found. Using zeros as starting parameters.
  warn('Non-invertible starting MA parameters found.')

```

SARIMAX Results

```
=====
Dep. Variable:          log_CPINSA4    No. Observations:          189
Model:                ARIMA(0, 0, 10)  Log Likelihood              768.413
Date:                 Tue, 15 Jul 2025  AIC                          -1512.826
Time:                 21:30:16         BIC                          -1473.925
Sample:               03-31-1961       HQIC                         -1497.066
                        - 03-31-2008
=====
```

Covariance Type: opg

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const          0.0395        0.005       7.731      0.000        0.029        0.050
ma.L1          1.5961        0.055      28.853      0.000        1.488        1.704
ma.L2          1.9248        0.116      16.662      0.000        1.698        2.151
ma.L3          2.2935        0.169      13.555      0.000        1.962        2.625
ma.L4          1.9670        0.200       9.857      0.000        1.576        2.358
ma.L5          1.6215        0.219       7.396      0.000        1.192        2.051
ma.L6          1.5860        0.229       6.936      0.000        1.138        2.034
ma.L7          1.2069        0.219       5.523      0.000        0.779        1.635
ma.L8          0.7927        0.171       4.638      0.000        0.458        1.128
ma.L9          0.5756        0.123       4.665      0.000        0.334        0.817
ma.L10         0.1923        0.071       2.693      0.007        0.052        0.332
sigma2         1.671e-05    1.62e-06    10.303      0.000    1.35e-05    1.99e-05
=====
```

```
===
Ljung-Box (L1) (Q):                0.00    Jarque-Bera (JB):
156.44
Prob(Q):                          0.99    Prob(JB):
0.00
Heteroskedasticity (H):            0.16    Skew:
1.17
Prob(H) (two-sided):              0.00    Kurtosis:
6.80
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
/home/ouslan/Documents/Github/ECON-124/.venv/lib/python3.10/site-
packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood
optimization failed to converge. Check mle_retvals
  warnings.warn("Maximum Likelihood optimization failed to "
```

```
[32]: arma_mod = ARIMA(df['log_CPINSA4'].dropna(), order=(5, 0, 10)).fit()
      print(arma_mod.summary())
```

SARIMAX Results

```

=====
Dep. Variable:          log_CPINSA4    No. Observations:          189
Model:                ARIMA(5, 0, 10)  Log Likelihood              791.966
Date:                 Tue, 15 Jul 2025  AIC                           -1549.932
Time:                 21:30:16          BIC                           -1494.822
Sample:               03-31-1961       HQIC                          -1527.605
                    - 03-31-2008

```

Covariance Type: opg

```

=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
const          0.0375     0.012     3.003     0.003     0.013     0.062
ar.L1          1.5931     0.567     2.808     0.005     0.481     2.705
ar.L2         -0.5906     0.845    -0.699     0.485    -2.247     1.066
ar.L3          0.0890     0.648     0.137     0.891    -1.180     1.358
ar.L4         -0.4775     0.718    -0.666     0.506    -1.884     0.929
ar.L5          0.3435     0.348     0.987     0.324    -0.338     1.025
ma.L1          0.0075     0.576     0.013     0.990    -1.121     1.136
ma.L2          0.0831     0.409     0.203     0.839    -0.719     0.886
ma.L3          0.1381     0.431     0.321     0.748    -0.706     0.982
ma.L4         -0.4304     0.456    -0.944     0.345    -1.324     0.463
ma.L5          0.3823     0.285     1.339     0.181    -0.177     0.942
ma.L6          0.3555     0.312     1.141     0.254    -0.255     0.966
ma.L7          0.1501     0.403     0.372     0.710    -0.641     0.941
ma.L8         -0.0303     0.186    -0.163     0.871    -0.395     0.334
ma.L9          0.0529     0.143     0.371     0.711    -0.227     0.333
ma.L10         -0.1239     0.170    -0.730     0.465    -0.457     0.209
sigma2         1.277e-05   9.72e-07   13.134     0.000   1.09e-05   1.47e-05
=====

```

```

===
Ljung-Box (L1) (Q):          0.05   Jarque-Bera (JB):
117.79
Prob(Q):                    0.82   Prob(JB):
0.00
Heteroskedasticity (H):      0.16   Skew:
0.14
Prob(H) (two-sided):         0.00   Kurtosis:
6.86
=====
===

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


```
[33]: arma_mod = ARIMA(df['log_CPINSA4'].dropna(), order=(6, 0, 7)).fit()
print(arma_mod.summary())
```

SARIMAX Results

```
=====
Dep. Variable:          log_CPINSA4    No. Observations:          189
Model:                ARIMA(6, 0, 7)    Log Likelihood              791.248
Date:                 Tue, 15 Jul 2025    AIC                       -1552.496
Time:                  21:30:17          BIC                       -1503.870
Sample:                03-31-1961        HQIC                      -1532.796
                  - 03-31-2008
```

Covariance Type: opg

```
=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
const          0.0383     0.011     3.337     0.001     0.016     0.061
ar.L1          1.3653     0.955     1.429     0.153    -0.507     3.238
ar.L2         -0.3581     1.621    -0.221     0.825    -3.536     2.820
ar.L3          0.3838     1.052     0.365     0.715    -1.678     2.446
ar.L4         -0.7411     0.380    -1.949     0.051    -1.486     0.004
ar.L5          0.2318     0.414     0.560     0.576    -0.580     1.044
ar.L6          0.0712     0.318     0.224     0.823    -0.553     0.695
ma.L1          0.2100     0.970     0.216     0.829    -1.692     2.112
ma.L2          0.2335     0.854     0.273     0.785    -1.441     1.908
ma.L3         -0.0672     0.430    -0.156     0.876    -0.909     0.775
ma.L4         -0.4725     0.371    -1.275     0.202    -1.199     0.254
ma.L5          0.1967     0.654     0.301     0.764    -1.085     1.478
ma.L6          0.1879     0.584     0.322     0.747    -0.956     1.332
ma.L7          0.2809     0.280     1.004     0.315    -0.268     0.829
sigma2         1.304e-05  1.07e-06  12.198     0.000    1.09e-05  1.51e-05
=====
```

```
===
Ljung-Box (L1) (Q):                0.04    Jarque-Bera (JB):
133.37
Prob(Q):                            0.84    Prob(JB):
0.00
Heteroskedasticity (H):              0.14    Skew:
0.13
Prob(H) (two-sided):                 0.00    Kurtosis:
7.11
=====
===
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

/home/ouslan/Documents/Github/ECON-124/.venv/lib/python3.10/site-

```
packages/statsmodels/base/model.py:607: ConvergenceWarning: Maximum Likelihood
optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to "
```

3.8 Problem H

- looking at the AIC the MA(10) seems to preform the best

3.9 Problem I

```
[34]: df["quarter"] = df.index.quarter
quarter_dummies = pd.get_dummies(df["quarter"], prefix="Q", drop_first=True).
      ↳ astype(int)
df = pd.concat([df, quarter_dummies], axis=1)
X = df[[col for col in df.columns if col.startswith("Q_")]]
X = sm.add_constant(X)
y = df["log_CPINSA4"]

model = sm.OLS(y, X, missing='drop').fit()
print(model.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          log_CPINSA4      R-squared:                0.000
Model:                  OLS              Adj. R-squared:           -0.016
Method:                 Least Squares    F-statistic:                0.005292
Date:                  Tue, 15 Jul 2025   Prob (F-statistic):         0.999
Time:                  21:30:17          Log-Likelihood:             428.53
No. Observations:      189              AIC:                       -849.1
Df Residuals:          185              BIC:                       -836.1
Df Model:               3
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0405	0.004	11.086	0.000	0.033	0.048
Q_2	0.0005	0.005	0.093	0.926	-0.010	0.011
Q_3	0.0006	0.005	0.108	0.914	-0.010	0.011
Q_4	0.0005	0.005	0.105	0.917	-0.010	0.011

```

=====
Omnibus:                42.124    Durbin-Watson:           0.045
Prob(Omnibus):           0.000    Jarque-Bera (JB):        62.900
Skew:                    1.275    Prob(JB):                2.19e-14
Kurtosis:                4.218    Cond. No.                4.76
=====

```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly

specified.

3.10 Problem J


```
[35]: df["residuals"] = model.resid
plt.figure(figsize=(10, 5))
plt.plot(df.index.to_timestamp(), df["residuals"], marker='o', linestyle='-')
plt.axhline(0, color='red', linestyle='--')
plt.title("Regression Residuals: log(CPinsa_t / CPinsa_{t-1}) ~ Quarterly_␣
↪Dummies")
plt.xlabel("Date")
plt.ylabel("Residual")
plt.grid(True)
plt.tight_layout()
plt.show()
```




- the residuals appear to be stationary

3.11 Problem K

```
[36]: fig = plot_acf(df['log_CPinsa4'].dropna(), lags=25)
plt.show()
fig = plot_pacf(df['log_CPinsa4'].dropna(), lags=25)
plt.show()
```





- the residuals seem to be autocorrelated by round 10 periods

4 Problem

4.1 Problem A

```
[37]: df = pd.read_excel("data/QUARTERLY-1.xls")
df['date'] = pd.PeriodIndex(df['Date'], freq='Q')
df.set_index('date', inplace=True)
# df = df[["CPINSA", "Date"]]
df["s"] = df["r10"] - df["Tbill"]
df = df[["s", "Date"]]
df
```


```
[37]:
```

	s	Date
date		
1960Q1	0.54334	1960Q1
1960Q2	1.17000	1960Q2
1960Q3	1.44000	1960Q3
1960Q4	1.52667	1960Q4
1961Q1	1.41000	1961Q1

```
...      ...      ...
2007Q1 -0.30000 2007Q1
2007Q2  0.09000 2007Q2
2007Q3  0.33667 2007Q3
2007Q4  0.75000 2007Q4
2008Q1  1.51333 2008Q1
```

[193 rows x 2 columns]


```
[38]: plt.figure(figsize=(8, 5))
plt.plot(df.index.to_timestamp(), df['s'], marker='o')
plt.title('Quarterly Data')
plt.xlabel('Date')
plt.ylabel('TB3mo')
plt.grid(True)
plt.tight_layout()
plt.show()
```




- it looks stationary

4.2 Problem B

```
[39]: fig = plot_acf(df['s'], lags=25)
plt.show()
plot_pacf(df['s'], lags=25)
plt.show()
```





- it seems to be autocorrelated by 4 units and partial autocorrelated by 1

```
[40]: res_ar2 = AutoReg(df['s'].dropna(), lags =2).fit()
print(res_ar2.summary())
df["residuals"] = res_ar2.resid
plt.figure(figsize=(10, 5))
plt.plot(df.index.to_timestamp(), df["residuals"], marker='o', linestyle='-')
plt.axhline(0, color='red', linestyle='--')
plt.title("Regression Residuals: log(CPinsa_t / CPinsa_{t-1}) ~ Quarterly_
↳Dummies")
plt.xlabel("Date")
plt.ylabel("Residual")
plt.grid(True)
plt.tight_layout()
plt.show()
```

AutoReg Model Results


Dep. Variable:	s	No. Observations:	193
Model:	AutoReg(2)	Log Likelihood	-150.856
Method:	Conditional MLE	S.D. of innovations	0.533
Date:	Tue, 15 Jul 2025	AIC	309.713

Time: 21:30:17 BIC 322.722
Sample: 09-30-1960 HQIC 314.982
- 03-31-2008

	coef	std err	z	P> z	[0.025	0.975]
const	0.1892	0.059	3.185	0.001	0.073	0.306
s.L1	1.1086	0.070	15.775	0.000	0.971	1.246
s.L2	-0.2450	0.070	-3.488	0.000	-0.383	-0.107

Roots

	Real	Imaginary	Modulus	Frequency
AR.1	1.2440	+0.0000j	1.2440	0.0000
AR.2	3.2815	+0.0000j	3.2815	0.0000



```
[41]: residuals = res_ar2.resid

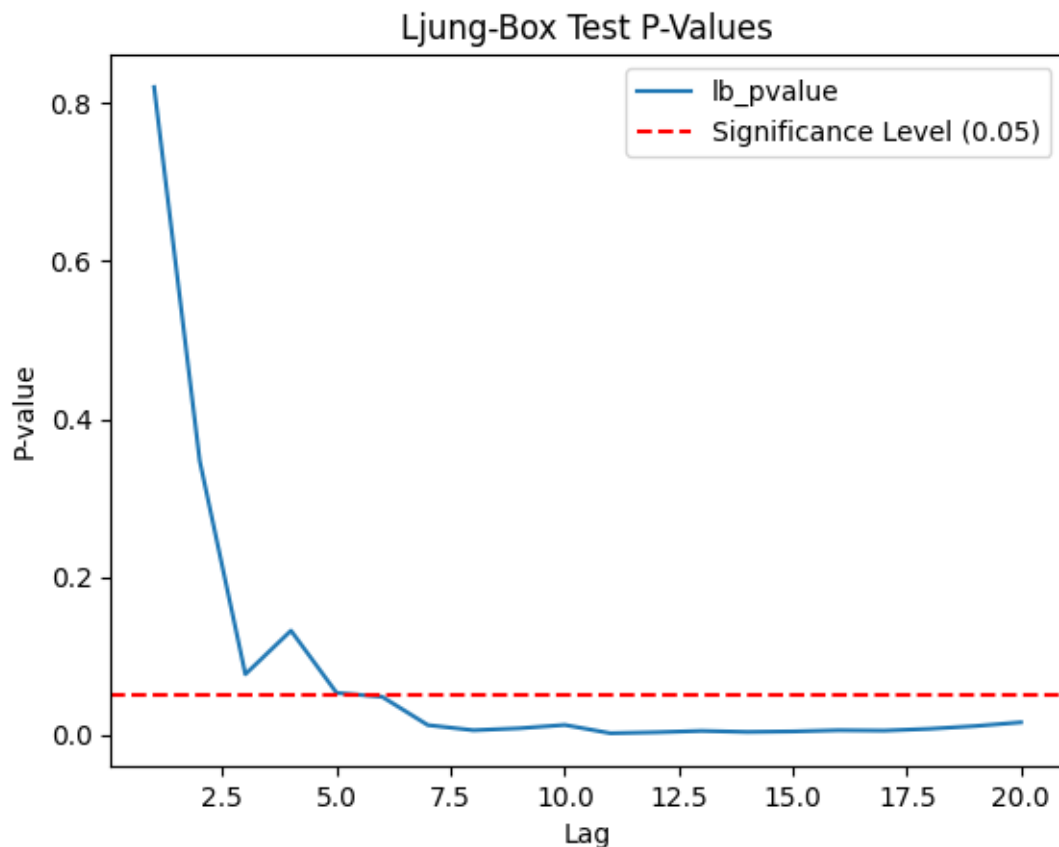
ljung_box_results = acorr_ljungbox(residuals, lags=range(1, 21), return_df=True)

print(ljung_box_results)

# Optional: Plot p-values to visualize autocorrelation
ljung_box_results['lb_pvalue'].plot(title='Ljung-Box Test P-Values')
plt.axhline(y=0.05, color='red', linestyle='--', label='Significance Level (0.05)')
plt.xlabel('Lag')
plt.ylabel('P-value')
```

```
plt.legend()  
plt.show()
```

	lb_stat	lb_pvalue
1	0.051333	0.820760
2	2.117369	0.346912
3	6.858840	0.076536
4	7.078838	0.131780
5	10.925209	0.052883
6	12.709362	0.047891
7	17.987541	0.012026
8	21.684678	0.005535
9	22.297068	0.007984
10	22.618691	0.012245
11	29.728266	0.001748
12	29.835833	0.002956
13	30.007268	0.004698
14	32.537869	0.003357
15	33.486681	0.004018
16	33.949568	0.005519
17	35.670142	0.005074
18	35.869956	0.007332
19	35.901840	0.010852
20	35.957500	0.015559



- there seem to be autocorroration for periods after 5

4.3 Problem E

```
[42]: res_ar7 = AutoReg(df['s'].dropna(), lags =7).fit()
      print(res_ar7.summary())
```

AutoReg Model Results						
=====						
Dep. Variable:	s		No. Observations:		193	
Model:	AutoReg(7)		Log Likelihood		-139.381	
Method:	Conditional MLE		S.D. of innovations		0.512	
Date:	Tue, 15 Jul 2025		AIC		296.763	
Time:	21:30:18		BIC		325.794	
Sample:	12-31-1961		HQIC		308.527	
	- 03-31-2008					
=====						
	coef	std err	z	P> z	[0.025	0.975]

const	0.2104	0.067	3.152	0.002	0.080	0.341
s.L1	1.1768	0.073	16.192	0.000	1.034	1.319

s.L2	-0.4658	0.109	-4.258	0.000	-0.680	-0.251
s.L3	0.3861	0.112	3.443	0.001	0.166	0.606
s.L4	-0.3386	0.113	-2.996	0.003	-0.560	-0.117
s.L5	0.3188	0.112	2.840	0.005	0.099	0.539
s.L6	-0.3791	0.109	-3.466	0.001	-0.593	-0.165
s.L7	0.1504	0.073	2.066	0.039	0.008	0.293

Roots

	Real	Imaginary	Modulus	Frequency
AR.1	-0.9208	-0.8352j	1.2432	-0.3828
AR.2	-0.9208	+0.8352j	1.2432	0.3828
AR.3	0.1872	-1.2949j	1.3083	-0.2271
AR.4	0.1872	+1.2949j	1.3083	0.2271
AR.5	1.2281	-0.3642j	1.2810	-0.0459
AR.6	1.2281	+0.3642j	1.2810	0.0459
AR.7	1.5317	-0.0000j	1.5317	-0.0000

4.4 Problem F

```
[43]: residuals = res_ar7.resid


ljung_box_results = acorr_ljungbox(residuals, lags=range(1, 21), return_df=True)

print(ljung_box_results)

ljung_box_results['lb_pvalue'].plot(title='Ljung-Box Test P-Values')
plt.axhline(y=0.05, color='red', linestyle='--', label='Significance Level (0.05)')
plt.xlabel('Lag')
plt.ylabel('P-value')
plt.legend()
plt.show()
```

	lb_stat	lb_pvalue
1	0.000031	0.995573
2	0.107644	0.947601
3	0.145260	0.985901
4	0.228593	0.993945
5	0.681262	0.983981
6	0.813668	0.991702
7	1.444265	0.984177
8	4.386175	0.820710
9	5.651757	0.774202
10	5.669816	0.842199
11	12.263550	0.344163

12	12.683548	0.392455
13	12.855834	0.459008
14	13.448330	0.491565
15	14.105729	0.517526
16	14.532605	0.559100
17	16.863113	0.463677
18	16.924938	0.528271
19	17.146294	0.579957
20	18.076740	0.582353



- there does not seem to be autocorrelation

4.5 Problem G

```
[44]: print(f"AIC AR(2): {res_ar2.aic}")
      print(f"AIC AR(7): {res_ar7.aic}")

      print(f"BIC AR(2): {res_ar2.bic}")
      print(f"BIC AR(7): {res_ar7.bic}")
```

AIC AR(2): 309.7126128947031

```
AIC AR(7): 296.76261348236477
BIC AR(2): 322.72170660688965
BIC AR(7): 325.7943335457836
```

4.6 Problem H

```
[45]: df2 = df.head(-10)
      y = df2['s'].dropna()

      forecast_ar2 = res_ar2.predict(start=res_ar2.model._hold_back, end=len(y)-1)

      error_ar2 = y[res_ar2.model._hold_back:] - forecast_ar2
      error_ar2
```

```
[45]: date
      1960Q3    0.086837
      1960Q4    0.027696
      1961Q1   -0.118914
      1961Q2    0.084988
      1961Q3    0.190625
      ...
      2004Q3   -0.521260
      2004Q4   -0.272081
      2005Q1   -0.143164
      2005Q2   -0.322460
      2005Q3   -0.347333
      Freq: Q-DEC, Length: 181, dtype: float64
```

```
[46]: y = df2['s'].dropna()

      res_ar7 = AutoReg(df2['s'].dropna(), lags=7).fit()

      forecast_ar7 = res_ar7.predict(start=res_ar7.model._hold_back, end=len(y)-1)

      error_ar7 = y[res_ar7.model._hold_back:] - forecast_ar7
      error_ar7
```

```
[46]: date
      1961Q4   -0.096905
      1962Q1   -0.124619
      1962Q2   -0.089512
      1962Q3   -0.030522
      1962Q4   -0.080839
      ...
      2004Q3   -0.449768
      2004Q4   -0.268528
      2005Q1   -0.218167
```

```
2005Q2    -0.097570
2005Q3    -0.433046
Freq: Q-DEC, Length: 176, dtype: float64
```

```
[47]: mse_ar2 = (error_ar2**2).mean()
      mse_ar7 = (error_ar7**2).mean()

      print(f'MSE AR(2): {mse_ar2:.4f}')
      print(f'MSE AR(7): {mse_ar7:.4f}')
```

```
MSE AR(2): 0.2933
```

```
MSE AR(7): 0.2697
```

- The AR(7) seems to have a smaller forecast error than AR(2)

4.7 Problem I

```
[48]: forecast_ar2 = res_ar2.predict(start=len(y), end=len(y)+9)

      # Forecast error
      error_ar2 = df['s'].tail(10) - forecast_ar2
      error_ar2
```

```
[48]: date
      2005Q4    -0.144416
      2006Q1    -0.518433
      2006Q2     0.124451
      2006Q3    -0.574535
      2006Q4    -0.336027
      2007Q1    -0.202987
      2007Q2     0.168873
      2007Q3    -0.025795
      2007Q4     0.209614
      2008Q1     0.575152
      Freq: Q-DEC, dtype: float64
```

```
[49]: forecast_ar7 = res_ar7.predict(start=len(y), end=len(y)+9)

      # Forecast error
      error_ar7 = df['s'].tail(10) - forecast_ar7
      forecast_ar7
```

```
[49]: 2005Q4    0.585641
      2006Q1    0.592219
      2006Q2    0.707072
      2006Q3    0.791292
      2006Q4    0.896282
      2007Q1    1.037073
```

```
2007Q2    1.173433
2007Q3    1.272218
2007Q4    1.326425
2008Q1    1.368948
Freq: Q-DEC, dtype: float64
```

```
[50]: mse_ar2 = (error_ar2**2).mean()
      mse_ar7 = (error_ar7**2).mean()

      print(f'MSE AR(2): {mse_ar2:.4f}')
      print(f'MSE AR(7): {mse_ar7:.4f}')
```

```
MSE AR(2): 0.1193
```

```
MSE AR(7): 0.6474
```

- the AR(2) seems to perform better than the AR(7) at forecasting the 10 steps. This is surprising given that looking at the entire series the AR(7) fits better the historical data