ECON 124: Midterm #2

Due on Jul 9, 2025

Dr. Deniz Baglan

Alejandro Ouslan

Problem 1

Use the data in **FERTIL2.XLSX** to answer this question.

(a) Estimate the model

$$children = \beta_0 + \beta_1 age + \beta_2 age^2 \beta_3 educ + \beta_4 electricity + \beta_5 urban + \epsilon$$

And report the usual and heteroskedasticity-robust standard errors. Are the robust standard errors always bigger than the non robust ones?

Answer: It seems that the robust standard errors are generally larger than the non robust ones, but not neccesarily always the case.

Table 1: OLS Regression Results: Children ~Age, Age², Education, Electricity, Urban

Variable	**Coef.**	**Std. Err.**	**t**	**P>** $\lvert t\rvert$	**[0.025**	**0.975]**
const	-4.2225	0.240	-17.580	0.000	-4.693	-3.752
age	0.3409	0.017	20.652	0.000	0.309	0.373
age^2	-0.0027	0.000	-10.086	0.000	-0.003	-0.002
education	-0.0752	0.006	-11.948	0.000	-0.088	-0.063
electricity	-0.3100	0.069	-4.493	0.000	-0.445	-0.175
urban	-0.2000	0.047	-4.301	0.000	-0.291	-0.109
Model statistics:						
R-squared		0.573				
Adj. R-squared		0.573				
F-statistic		1170	(Prob F-statistic = 0.000)			
No. Observations		4358				
Df Residuals		4352				
Df Model		5				
Log-Likelihood		-7806.3				
AIC		1.562e+04				
BIC		1.566e+04				
Durbin-Watson: 1.883						
Omnibus: 203.155, Prob(Omnibus): 0.000						
Jarque-Bera (JB): 715.135, Prob(JB): 5.13×10^{-156}						
Skew: 0.014, Kurtosis: 4.984						
Cond. No.: 1.07e+04						

Problem 1 continued on next page. . .

Table 2: OLS Regression Results: Children ∼Age, Age², Education, Electricity, Urban

Variable	**Coef.**	**Std. Err.**	**z**	**P>** $\lvert z\rvert$	**[0.025**	**0.975]**
const	-4.2225	0.244	-17.316	0.000	-4.700	-3.745
age	0.3409	0.019	17.780	0.000	0.303	0.379
age²	-0.0027	0.000	-7.821	0.000	-0.003	-0.002
education	-0.0752	0.006	-11.927	0.000	-0.088	-0.063
electricity	-0.3100	0.064	-4.848	0.000	-0.435	-0.185
urban	-0.2000	0.045	-4.399	0.000	-0.289	-0.111
Model statistics:						
R-squared		0.573				
Adj. R-squared		0.573				
F-statistic		1161	(Prob F-statistic = 0.000)			
No. Observations		4358				
Df Residuals		4352				
Df Model		5				
Log-Likelihood		-7806.3				
AIC		1.562e+04				
BIC		1.566e+04				
Durbin-Watson: 1.883						
Omnibus: 203.155, Prob(Omnibus): 0.000						
Jarque-Bera (JB): 715.135, Prob(JB): 5.13×10^{-156}						
Skew: 0.014, Kurtosis: 4.984						
Cond. No.: 1.07e+04						

Python Code

```python
import polars as pl
import statsmodels.formula.api as smf

# Question
# 1a
df = pl.read_excel("data/fertil2.xlsx")
df = df.select(
    pl.col(
        [
            "children",
            "age",
            "educ",
            "electric",
            "urban",
            "spirit",
            "protest",
            "catholic",
        ]
    )
)
df = df.with_columns(age2=pl.col("age") ** 2)
df = df.to_pandas()
```

Problem 1 continued on next page. . .

```python
model = smf.ols("children ~ age + age2 + educ + electric + urban", data=df).fit()
print(model.summary())

model = smf.ols("children ~ age + age2 + educ + electric + urban", data=df).fit(
    cov_type="HC1"
)
print(model.summary())
```

(b) Add the three religious dummy variables and test whether they are jointly significant. What are the p-values for the nonrobust and robust tests?
Answer: The p-values for the non-robust test is 0.0864, while the p-value for the robust test is 0.0911. It seems that robust tests are less likely to report something is significant, especially assuming standard errors are greater than non-robust ones.

Table 3: OLS Regression Results: Children ∼ Age, Age², Education, Electricity, Urban, Spirit, Protest, Catholic

Variable	**Coef.**	**Std. Err.**	**t**	**P>** $\lvert t \rvert$	**[0.025**	**0.975]**
Intercept	-4.3147	0.243	-17.731	0.000	-4.792	-3.838
age	0.3419	0.017	20.696	0.000	0.309	0.374
age²	-0.0028	0.000	-10.139	0.000	-0.003	-0.002
education	-0.0762	0.006	-11.796	0.000	-0.089	-0.064
electricity	-0.3057	0.069	-4.429	0.000	-0.441	-0.170
urban	-0.2034	0.047	-4.366	0.000	-0.295	-0.112
spirit	0.1405	0.056	2.517	0.012	0.031	0.250
protest	0.0754	0.065	1.156	0.248	-0.052	0.203
catholic	0.1174	0.083	1.407	0.160	-0.046	0.281
Model statistics:						
R-squared		0.574				
Adj. R-squared		0.573				
F-statistic		732.6	(Prob F-statistic = 0.000)			
No. Observations		4358				
Df Residuals		4349				
Df Model		8				
Log-Likelihood		-7803.0				
AIC		1.562e+04				
BIC		1.568e+04				
Durbin-Watson: 1.887						
Omnibus: 202.228, Prob(Omnibus): 0.000						
Jarque-Bera (JB): 709.030, Prob(JB): 1.09×10^{-154}						
Skew: 0.016, Kurtosis: 4.976						
Cond. No.: 1.09e+04						

Problem 1 continued on next page. . .

Table 4: F-test Results

Statistic	**Value**	**Notes**
F-statistic	2.196	-
p-value	0.0864	-
Degrees of Freedom (denominator)	4350	-
Degrees of Freedom (numerator)	3	-

Table 5: OLS Regression Results: Children $\sim$ Age, Age², Education, Electricity, Urban, Spirit, Protest, Catholic (Robust Standard Errors)

Variable	**Coef.**	**Std. Err.**	**z**	**P$>|z|$**	**[0.025**	**0.975]**
Intercept	-4.3147	0.248	-17.389	0.000	-4.801	-3.828
age	0.3419	0.019	17.807	0.000	0.304	0.379
age²	-0.0028	0.000	-7.861	0.000	-0.003	-0.002
education	-0.0762	0.006	-11.860	0.000	-0.089	-0.064
electricity	-0.3057	0.064	-4.772	0.000	-0.431	-0.180
urban	-0.2034	0.046	-4.456	0.000	-0.293	-0.114
spirit	0.1405	0.056	2.487	0.013	0.030	0.251
protest	0.0754	0.066	1.140	0.254	-0.054	0.205
catholic	0.1174	0.079	1.483	0.138	-0.038	0.272
Model statistics:						
R-squared		0.574				
Adj. R-squared		0.573				
F-statistic		727.9	(Prob F-statistic = 0.000)			
No. Observations		4358				
Df Residuals		4349				
Df Model		8				
Log-Likelihood		-7803.0				
AIC		1.562e+04				
BIC		1.568e+04				
Durbin-Watson: 1.887						
Omnibus: 202.228, Prob(Omnibus): 0.000						
Jarque-Bera (JB): 709.030, Prob(JB): 1.09×10^{-154}						
Skew: 0.016, Kurtosis: 4.976						
Cond. No.: 1.09e+04						

Table 6: F-test Results

Statistic	**Value**	**Notes**
F-statistic	2.156	-
p-value	0.0911	-
Degrees of Freedom (denominator)	4350	-
Degrees of Freedom (numerator)	3	-

Problem 1 continued on next page. . .

Python Code

```
# assumes that the code in 1a ran
model = smf.ols(
    "children ~ age + age2 + educ + electric + urban + spirit + protest + catholic",
    data=df,
).fit()
print(model.summary())
print(model.f_test("spirit = protest = catholic = 0"))

model = smf.ols(
    "children ~ age + age2 + educ + electric + urban + spirit + protest + catholic",
    data=df,
).fit(cov_type="HC1")
print(model.summary())
print(model.f_test("spirit = protest = catholic = 0"))
```

(c) From the regression in part (b), obtain the fitted values $\hat{y}$ and the residuals, $\hat{\epsilon}$. Regress $\hat{\epsilon}^2 \sim \hat{y}$, and $\hat{\epsilon}^2 \sim \hat{y}^2$ and test the joint significance of the two regressors.

Table 7: OLS Regression Results: $\hat{u}^2 \sim \hat{y} + \hat{y}^2$

Variable	**Coef.**	**Std. Err.**	**t**	**P**$>|t|$	**[0.025**	**0.975]**
Intercept	0.3126	0.111	2.807	0.005	0.094	0.531
$\hat{y}$	-0.1489	0.102	-1.462	0.144	-0.348	0.051
$\hat{y}^2$	0.2668	0.020	13.607	0.000	0.228	0.305
Model statistics:						
R-squared		0.250				
Adj. R-squared		0.250				
F-statistic		726.1	(Prob F-statistic = 7.19e-273)			
No. Observations		4358				
Df Residuals		4355				
Df Model		2				
Log-Likelihood		-11803				
AIC		2.361e+04				
BIC		2.363e+04				
Durbin-Watson: 1.947						
Omnibus: 3446.975, Prob(Omnibus): 0.000						
Jarque-Bera (JB): 119444.435, Prob(JB): 0.000						
Skew: 3.503, Kurtosis: 27.672						
Cond. No.: 31.4						

Table 8: F-test Results

Statistic	**Value**	**Notes**
F-statistic	726.11	-
p-value	7.19×10^{-273}	-
Degrees of Freedom (denominator)	4358	-
Degrees of Freedom (numerator)	2	-

Python Code

```python
# assums that the code in problem 1 a ran
df["yhat"] = model.fittedvalues
df["u_hat"] = model.resid
df["u_hat2"] = df["u_hat"] ** 2
df["yhat2"] = df["yhat"] ** 2

model = smf.ols("u_hat2 ~ yhat + yhat2", data=df).fit()
print(model.summary())
print(model.f_test("yhat = yhat2 = 0"))
```

Problem 2

Use the data set **Movies**

Does viewing a violent movie lead to violent behavior? If so, the incidence of violent crimes, such as assaults, should rise following the release of a violent movie that attracts many viewers. Alternatively, movie viewing may substitute for other activities (such as alcohol consumption) that lead to violent behavior, so that assaults should fall when more viewers are attracted to the cinema. The dataset includes weekend U.S. attendance for strongly violent movies (such as Hannibal), mildly violent movies (such as Spider-Man), and nonviolent movies (such as Finding Nemo). The dataset also includes a count of the number of assaults for the same weekend in a subset of counties in the United States. Finally, the dataset includes indicators for year, month, whether the weekend is a holiday, and various measures of the weather.

(a) (i) Regress the logarithm of the number of assaults [ln_(assaults) = ln(assaults)] on the year and month indicators. Is there evidence of seasonality in assaults? That is, do there tend to be more assaults in some months than others? Explain.

Answer: In comparison to January there seems to be more assaults during late spring and early fall, especially in the summer. (May to Sepetember) So there do seem to be some seasonality as assaults are lower during winter months.

Problem 2 continued on next page...

Table 9: OLS Regression Results: ln(assaults)

Variable	**Coef.**	**Std. Err.**	**t**	**P**$>|t|$	**[0.025**	**0.975]**
Intercept	6.7276	0.012	561.447	0.000	6.704	6.751
year2	0.6949	0.012	60.043	0.000	0.672	0.718
year3	1.0084	0.012	87.090	0.000	0.986	1.031
year4	1.2454	0.012	107.557	0.000	1.223	1.268
year5	1.4145	0.012	122.757	0.000	1.392	1.437
year6	1.6953	0.012	146.577	0.000	1.673	1.718
year7	1.8509	0.012	159.962	0.000	1.828	1.874
year8	1.9013	0.012	164.285	0.000	1.879	1.924
year9	1.9437	0.012	167.868	0.000	1.921	1.966
year10	2.0702	0.012	175.002	0.000	2.047	2.093
month2	0.0169	0.013	1.297	0.195	-0.009	0.042
month3	0.0799	0.013	6.379	0.000	0.055	0.105
month4	0.1297	0.013	10.245	0.000	0.105	0.155
month5	0.1802	0.012	14.480	0.000	0.156	0.205
month6	0.1663	0.013	13.200	0.000	0.142	0.191
month7	0.1739	0.013	13.815	0.000	0.149	0.199
month8	0.1768	0.012	14.200	0.000	0.152	0.201
month9	0.1977	0.013	15.599	0.000	0.173	0.223
month10	0.1417	0.012	11.390	0.000	0.117	0.166
month11	0.0248	0.013	1.971	0.049	0.00007	0.050
month12	-0.0054	0.013	-0.429	0.668	-0.030	0.019
Model statistics:						
R-squared		0.992				
Adj. R-squared		0.991				
F-statistic		2948	(Prob F-statistic = 0.000)			
No. Observations		516				
Df Residuals		495				
Df Model		20				
Log-Likelihood		741.74				
AIC		-1441				
BIC		-1352				
Durbin-Watson: 1.997						
Omnibus: 147.335, Prob(Omnibus): 0.000						
Jarque-Bera (JB): 1613.470, Prob(JB): 0.000						
Skew: 0.907, Kurtosis: 11.471						
Cond. No.: 13.5						

Problem 2 continued on next page...

Python Code

```python
import polars as pl
import numpy as np
import statsmodels.formula.api as smf
from linearmodels.iv import IV2SLS

df = pl.read_excel("data/Movies.xlsx")
df = df.select(
    pl.all().exclude("month1", "year1"),
    ln_assaults=pl.col("assaults").log(),
    attendance=pl.col("attend_v") + pl.col("attend_m") + pl.col("attend_n"),
)
data = df.to_pandas()

variables = df.select(
    pl.all().exclude(
        "assaults", "ln_assaults", "^atten.*$", "^h_.*$", "^pr_.*$", "^w.*$"
    )
).columns
formula = "ln_assaults ~ " + " + ".join(variables)


model = smf.ols(formula=formula, data=df).fit()
print(model.summary())
```

(ii) Regress total movie attendance (attend = attend_v + attend_m + attend_n) on the year and month indicators. Is there evidence of seasonality in movie attendance? Explain.
Answer: In comparison to january there seems to be more movie attendance during the summer, especially in June, July, and oddly November. One could argue the re is seasonality, but November has an odd peak in movie attendance, which makes me think maybe attendance goes along with another variable like relea se date for movies.

Table 10: OLS Regression Results: attendance

Variable	**Coef.**	**Std. Err.**	**t**	**P**$>\|t\|$	**[0.025**	**0.975]**
Intercept	16.0396	0.737	21.770	0.000	14.592	17.487
year2	0.8972	0.712	1.261	0.208	-0.501	2.295
year3	2.0311	0.712	2.853	0.005	0.632	3.430
year4	3.3991	0.712	4.774	0.000	2.000	4.798
year5	2.9515	0.709	4.166	0.000	1.559	4.344
year6	2.5120	0.711	3.532	0.000	1.115	3.909
year7	3.1224	0.711	4.389	0.000	1.725	4.520
year8	5.0412	0.712	7.084	0.000	3.643	6.439
year9	4.4057	0.712	6.188	0.000	3.007	5.805
year10	3.7276	0.727	5.125	0.000	2.298	5.157
month2	-0.6613	0.800	-0.827	0.409	-2.233	0.911
month3	-2.1685	0.771	-2.814	0.005	-3.683	-0.654
month4	-3.4033	0.779	-4.371	0.000	-4.933	-1.873
month5	0.9224	0.765	1.205	0.229	-0.581	2.426
month6	2.9167	0.775	3.765	0.000	1.395	4.439
month7	5.3937	0.774	6.969	0.000	3.873	6.914
month8	1.2123	0.766	1.584	0.114	-0.292	2.716
month9	-5.4252	0.779	-6.962	0.000	-6.956	-3.894
month10	-3.3259	0.765	-4.347	0.000	-4.829	-1.823
month11	3.8795	0.774	5.009	0.000	2.358	5.401
month12	0.6774	0.779	0.869	0.385	-0.854	2.209
Model statistics:						
R-squared		0.480				
Adj. R-squared		0.459				
F-statistic		22.86	(Prob F-statistic = 7.85e-58)			
No. Observations		516				
Df Residuals		495				
Df Model		20				
Log-Likelihood		-1383.6				
AIC		2809				
BIC		2898				
Durbin-Watson: 1.439						
Omnibus: 69.939, Prob(Omnibus): 0.000						
Jarque-Bera (JB): 126.992, Prob(JB): 2.65e-28						
Skew: 0.808, Kurtosis: 4.816						
Cond. No.: 13.5						

Python Code

```
import polars as pl
import numpy as np
import statsmodels.formula.api as smf
from linearmodels.iv import IV2SLS

df = pl.read_excel("data/Movies.xlsx")
df = df.select(
    pl.all().exclude("month1", "year1"),
    ln_assaults=pl.col("assaults").log(),
    attendance=pl.col("attend_v") + pl.col("attend_m") + pl.col("attend_n"),
)
data = df.to_pandas()

variables = df.select(
    pl.all().exclude(
        "assaults", "ln_assaults", "^atten.*$", "^h_.*$", "^pr_.*$", "^w.*$"
    )
).columns
formula = "ln_assaults ~ " + " + ".join(variables)


model = smf.ols(formula=formula, data=df).fit()
print(model.summary())
```

(b) Regress ln_assaults on attend_v, attend_m, attend_n, the year and month indicators, and the weather and holiday control variables available in the data set

(i) Based on the regression, does viewing a strongly violent movie increase or decrease assaults? By how much? Is the estimated effect statistically significant?
Answer: When taking into account all the controls, it seems that viewing a strongly violent movie decreases assaults by 0.3 percent, which is significant since it's associated p-value is below 0.01.

Table 11: OLS Regression Results: ln_assaults

Variable	**Coef.**	**Std. Err.**	**t**	**P>** $\lvert t \rvert$	**[0.025**	**0.975]**
Intercept	6.9143	0.017	403.245	0.000	6.881	6.948
attend_v	-0.0032	0.001	-3.171	0.002	-0.005	-0.001
attend_m	-0.0031	0.001	-4.742	0.000	-0.004	-0.002
attend_n	-0.0021	0.001	-3.194	0.001	-0.003	-0.001
h_chris	-0.0879	0.023	-3.744	0.000	-0.134	-0.042
h_newyr	0.2453	0.023	10.700	0.000	0.200	0.290
h_easter	-0.0369	0.015	-2.537	0.012	-0.066	-0.008
h_july4	0.0352	0.020	1.736	0.083	-0.005	0.075
h_mem	0.0059	0.015	0.397	0.691	-0.023	0.035
h_labor	0.0241	0.014	1.697	0.090	-0.004	0.052
w_maxa	0.1099	0.013	8.155	0.000	0.083	0.136
w_maxb	0.1107	0.019	5.971	0.000	0.074	0.147
w_maxc	0.0423	0.070	0.607	0.544	-0.095	0.179
w_mina	-0.3405	0.040	-8.599	0.000	-0.418	-0.263
w_minb	-0.1725	0.027	-6.394	0.000	-0.226	-0.120
w_minc	-0.1196	0.017	-7.126	0.000	-0.153	-0.087
w_rain	-0.0323	0.013	-2.518	0.012	-0.057	-0.007
w_snow	-0.0612	0.030	-2.057	0.040	-0.120	-0.003

Python Code

```python
# aasums the code in 2 A i ran
variables = df.select(
    pl.all().exclude("assaults", "ln_assaults", "wkd_ind", "^pr_.*$", "^atten.*$")
).columns
formula = "ln_assaults ~ " + "attend_v + attend_m + attend_n + " + " + ".join(variables)

model = smf.ols(formula=formula, data=df).fit()
print(model.summary())
```

Problem 2 continued on next page. . .

Table 12: OLS Regression Results: ln_assaults

Variable	**Coef.**	**Std. Err.**	**t**	**P>** $\lvert t\rvert$	**[0.025**	**0.975]**
year2	0.7008	0.009	81.931	0.000	0.684	0.718
year3	1.0171	0.009	116.936	0.000	1.000	1.034
year4	1.2267	0.009	138.859	0.000	1.209	1.244
year5	1.3891	0.009	159.693	0.000	1.372	1.406
year6	1.6883	0.009	196.902	0.000	1.671	1.705
year7	1.8395	0.009	207.295	0.000	1.822	1.857
year8	1.8984	0.009	206.170	0.000	1.880	1.916
year9	1.9501	0.009	221.146	0.000	1.933	1.967
year10	2.0721	0.009	229.747	0.000	2.054	2.090
month2	-0.0073	0.010	-0.759	0.448	-0.026	0.012
month3	0.0126	0.010	1.241	0.215	-0.007	0.032
month4	0.0019	0.012	0.150	0.881	-0.023	0.026
month5	0.0079	0.015	0.546	0.585	-0.021	0.037
month6	-0.0290	0.016	-1.852	0.065	-0.060	0.002
month7	-0.0343	0.018	-1.929	0.054	-0.069	0.001
month8	-0.0385	0.017	-2.280	0.023	-0.072	-0.005
month9	-0.0129	0.015	-0.846	0.398	-0.043	0.017
month10	-0.0025	0.013	-0.197	0.844	-0.027	0.022
month11	-0.0432	0.011	-4.006	0.000	-0.064	-0.022
month12	-0.0305	0.010	-3.084	0.002	-0.050	-0.011
Model statistics:						
R-squared		0.996				
Adj. R-squared		0.996				
F-statistic		3139.0	(Prob F-statistic = 0.00)			
No. Observations		516				
Df Residuals		478				
Df Model		37				
Log-Likelihood		924.66				
AIC		-1773.0				
BIC		-1612.0				
Durbin-Watson: 1.917						
Omnibus: 126.823, Prob(Omnibus): 0.000						
Jarque-Bera (JB): 1910.561, Prob(JB): 0.00						
Skew: -0.612, Kurtosis: 12.347						
Cond. No.: 474						

Problem 2 continued on next page...

(ii) Does attendance at strongly violent movies affect assaults differently than attendance at moderately violent movies? Differently than attendance at nonviolent movies?
Answer: after doing the test we get a p value of 0.000 thus we reject the null hypothesis of that no of the moves have an impact on assault and thus conlude there is significant evidence to conclude that there is an association between the movies and assaults.

Table 13: F-test Results

Statistic	**Value**	**Notes**
F-statistic	8.12	-
p-value	2.77×10^{-5}	-
Degrees of Freedom (denominator)	478	-
Degrees of Freedom (numerator)	3	-

Python Code

```python
# aasums the code in 2 B i ran
print(model.f_test("attend_v = attend_m = attend_n = 0"))
```

(iii) A strongly violent blockbuster movie is released, and the weekend's attendance at strongly violent movies increases by 6 million; meanwhile, attendance falls by 2 million for moderately violent movies and by 1 million for nonviolent movies. What is the predicted effect on assaults? Construct a 95% confidence interval for the change in assaults.

Table 14: Predicted Percentage Change in Assaults

Statistic	**Value**
Predicted % change in assaults	-1.06%
95% Confidence Interval (CI)	[-2.09%, -0.01%]

Python Code

```python
coeffs = model.params[["attend_v", "attend_m", "attend_n"]].values
cov = (
    model.cov_params()
    .loc[["attend_v", "attend_m", "attend_n"], ["attend_v", "attend_m", "attend_n"]]
    .values
)

delta_x = np.array([6, -2, -1])

delta_ln_assaults = np.dot(delta_x, coeffs)

std_error = np.sqrt(np.dot(delta_x, np.dot(cov, delta_x)))

lower = delta_ln_assaults - 1.96 * std_error
upper = delta_ln_assaults + 1.96 * std_error

percent_change = 100 * (np.exp(delta_ln_assaults) - 1)
ci_lower = 100 * (np.exp(lower) - 1)
```

```
ci_upper = 100 * (np.exp(upper) - 1)

print(f"Predicted % change in assaults: {percent_change:.2f}%")
print(f"95% CI: [{ci_lower:.2f}%, {ci_upper:.2f}%]")
```

(c) It is difficult to control for all the variables that affect assaults and that might be correlated with movie attendance. For example, the effect of the weather on assaults and movie attendance is only crudely approximated by the weather variables in the data set. However, the data set does include a set of instruments, pr_attend_v, pr_attend_m, and pr_attend_n, that are correlated with attendance but are (arguably) uncorrelated with weekend-specific factors (such as the weather) that affect both assaults and movie attendance. These instruments use historical attendance patterns, not information on a particular weekend, to predict a film's attendance in a given weekend. For example, if a film's attendance is high in the second week of its release, then this can be used to predict that its attendance was also high in the first week of its release. (The details of the construction of these instruments are available in the Dahl and DellaVigna's paper on Canvas) Run the regression from part (b) (including year, month, holiday, and weather controls) but now using pr_attend_v, pr_attend_m, and pr_attend_n as instruments for attend_v, attend_m, and attend_n. Use this regression to answer (b)(i)–(b)(iii).

(i) **Answer:** When taking into account all the controls, it seems that viewing a strongly violent movie decreases assaults by 9.6 percent, which is significant since it's associated p-value is below 0.00.

Table 15: IV-2SLS Estimation Results: ln_assaults

Variable	**Coef.**	**Std. Err.**	**t-stat**	**P-value**	**[0.025**	**0.975]**
attend_m	0.0958	0.0144	6.6340	0.0000	0.0675	0.1241
attend_n	0.1246	0.0132	9.4294	0.0000	0.0987	0.1505
h_chris	-1.0922	0.5642	-1.9359	0.0529	-2.1979	0.0136
h_newyr	-0.9118	0.5940	-1.5350	0.1248	-2.0761	0.2525
h_easter	-0.0891	0.1748	-0.5096	0.6103	-0.4316	0.2535
h_july4	-0.1047	0.2475	-0.4229	0.6724	-0.5898	0.3805
h_mem	-1.0002	0.2005	-4.9889	0.0000	-1.3932	-0.6073
h_labor	-0.0863	0.1583	-0.5454	0.5855	-0.3965	0.2239
w_maxa	0.6434	0.1841	3.4941	0.0005	0.2825	1.0043
w_maxb	0.3844	0.2536	1.5160	0.1295	-0.1126	0.8814
w_maxc	-1.2307	0.9379	-1.3121	0.1895	-3.0690	0.6076
w_mina	2.7672	0.8390	3.2980	0.0010	1.1227	4.4116
w_minb	3.6024	0.5503	6.5462	0.0000	2.5238	4.6810
w_minc	3.0501	0.3225	9.4574	0.0000	2.4180	3.6822
w_rain	1.7415	0.2269	7.6745	0.0000	1.2967	2.1862
w_snow	1.4631	0.6946	2.1063	0.0352	0.1016	2.8245
year2	1.5297	0.1526	10.022	0.0000	1.2306	1.8289
year3	1.7852	0.1677	10.646	0.0000	1.4565	2.1139
year4	2.2659	0.1728	13.117	0.0000	1.9273	2.6045
year5	2.4164	0.1830	13.203	0.0000	2.0577	2.7751
year6	2.3753	0.1693	14.033	0.0000	2.0436	2.7070
year7	2.6131	0.1711	15.277	0.0000	2.2779	2.9484
year8	2.3433	0.1845	12.699	0.0000	1.9816	2.7050
year9	2.3747	0.1789	13.276	0.0000	2.0241	2.7253
year10	2.6030	0.1753	14.847	0.0000	2.2594	2.9467
month2	1.3267	0.2045	6.4865	0.0000	0.9258	1.7275
month3	2.2704	0.2040	11.132	0.0000	1.8706	2.6701
month4	3.1781	0.1930	16.468	0.0000	2.7998	3.5563
month5	3.3692	0.2405	14.010	0.0000	2.8978	3.8405
month6	2.7740	0.2756	10.064	0.0000	2.2338	3.3143
month7	2.5672	0.2974	8.6318	0.0000	1.9843	3.1502
month8	3.0297	0.2606	11.625	0.0000	2.5189	3.5405
month9	3.8706	0.2021	19.151	0.0000	3.4744	4.2667
month10	3.4674	0.1796	19.302	0.0000	3.1153	3.8195
month11	1.6990	0.2377	7.1471	0.0000	1.2331	2.1650
month12	1.2190	0.2251	5.4149	0.0000	0.7778	1.6602
attend_v	0.1199	0.0182	6.5803	0.0000	0.0842	0.1556
Model statistics:						
R-squared		0.9918				
Adj. R-squared		0.9912				
F-statistic		1.08e+05	(Prob F-statistic = 0.0000)			
No. Observations		516				
P-value (F-stat)		0.0000				
Distribution		chi2(37)				
Cov. Estimator		robust				

Python Code

```
endog = data[['attend_v']]
other_controls  = df.select(pl.all().exclude("assaults", "ln_assaults", "wkd_ind", "^atten
exog = data[['attend_m', 'attend_n'] + other_controls]
instruments = data[["pr_attend_v", "pr_attend_m", "pr_attend_n"]]
ivolsmod = IV2SLS(dependent=data[["ln_assaults"]], endog=endog, exog=exog, instruments=ins
res_ivols = ivolsmod.fit()
print(res_ivols.summary)
```

(ii) **Answer:** after doing the test we get a p value of 0.000 thus we reject the null hypothesis of that no of the moves have an impact on assault and thus conlude there is significant evidence to conclude that there is an association between the movies and assaults.

Table 16: Linear Equality Hypothesis Test

Statistic	**Value**	**Notes**
Test Statistic	96.8238	-
p-value	0.0000	-
Distribution	chi2(3)	-

Python Code

```
print(res_ivols.wald_test(formula="attend_v = attend_m = attend_n = 0"))
```

Table 17: Predicted Percentage Change in Assaults

Statistic	**Value**
Predicted % change in assaults	49.64%
95% Confidence Interval (CI)	[24.24%, 80.24%]

Python Code

(iii)

```
#iii
coeffs = res_ivols.params[['attend_v', 'attend_m', 'attend_n']].values
cov = res_ivols.cov.loc[['attend_v', 'attend_m', 'attend_n'], ['attend_v', 'attend_m', 'at

delta_x = np.array([6, -2, -1])

delta_ln_assaults = np.dot(delta_x, coeffs)

std_error = np.sqrt(np.dot(delta_x, np.dot(cov, delta_x)))

lower = delta_ln_assaults - 1.96 * std_error
upper = delta_ln_assaults + 1.96 * std_error

percent_change = 100 * (np.exp(delta_ln_assaults) - 1)
ci_lower = 100 * (np.exp(lower) - 1)
ci_upper = 100 * (np.exp(upper) - 1)
```

```
print(f"Predicted % change in assaults: {percent_change:.2f}%")
print(f"95% CI: [{ci_lower:.2f}%, {ci_upper:.2f}%]")
```

(d) It is difficult to control for all the variables that affect assaults and that might be correlated with movie attendance. For example, the effect of the weather on assaults and movie attendance is only crudely approximated by the weather variables in the data set. However, the data set does include a set of instruments, pr_attend_v, pr_attend_m, and pr_attend_n, that are correlated with attendance but are (arguably) uncorrelated with weekend-specific factors (such as the weather) that affect both assaults and movie attendance. These instruments use historical attendance patterns, not information on a particular weekend, to predict a film's attendance in a given weekend. For example, if a film's attendance is high in the second week of its release, then this can be used to predict that its attendance was also high in the first week of its release. (The details of the construction of these instruments are available in the Dahl and DellaVigna's paper on Canvas) Run the regression from part (b) (including year, month, holiday, and weather controls) but now using pr_attend_v, pr_attend_m, and pr_attend_n as instruments for attend_v, attend_m, and attend_n. Use this regression to answer (b)(i)–(b)(iii).

(i) **Answer:** When taking into account all the controls, it seems that viewing a stro ngly violent movie decreases assaults by 12.22 percent, which is significant since it's associated p-value is below 0.00.

Table 18: IV-2SLS Estimation Results: ln_assaults

Variable	**Coef.**	**Std. Err.**	**t-stat**	**P-value**	**[0.025**	**0.975]**
attend_m	0.1222	0.0157	7.7982	0.0000	0.0915	0.1529
attend_n	0.1431	0.0141	10.184	0.0000	0.1156	0.1707
h_chris	-1.2236	0.6015	-2.0342	0.0419	-2.4025	-0.0447
h_newyr	-1.0918	0.6409	-1.7035	0.0885	-2.3479	0.1644
h_easter	-0.0730	0.1730	-0.4223	0.6728	-0.4120	0.2660
h_july4	-0.0104	0.1988	-0.0525	0.9581	-0.4000	0.3792
h_mem	-1.1828	0.2242	-5.2759	0.0000	-1.6222	-0.7434
h_labor	-0.0397	0.1749	-0.2272	0.8203	-0.3825	0.3031
w_maxa	0.5675	0.1919	2.9574	0.0031	0.1914	0.9436
w_maxb	0.2810	0.2805	1.0017	0.3165	-0.2688	0.8307
w_maxc	-1.1101	1.0709	-1.0366	0.2999	-3.2091	0.9889
w_mina	2.7290	0.8219	3.3203	0.0009	1.1181	4.3399
w_minb	3.3800	0.5602	6.0335	0.0000	2.2820	4.4779
w_minc	2.6254	0.3399	7.7231	0.0000	1.9591	3.2917
w_rain	1.5124	0.2337	6.4712	0.0000	1.0544	1.9705
w_snow	1.1629	0.6960	1.6708	0.0948	-0.2013	2.5271
year2	1.4104	0.1580	8.9280	0.0000	1.1008	1.7201
year3	1.5671	0.1710	9.1649	0.0000	1.2320	1.9023
year4	2.0745	0.1765	11.751	0.0000	1.7285	2.4206
year5	2.2178	0.1824	12.159	0.0000	1.8603	2.5753
year6	2.2026	0.1744	12.629	0.0000	1.8608	2.5444
year7	2.4496	0.1765	13.879	0.0000	2.1036	2.7955
year8	2.2256	0.1913	11.632	0.0000	1.8506	2.6006
year9	2.1518	0.1826	11.781	0.0000	1.7938	2.5098
year10	2.3896	0.1753	13.632	0.0000	2.0460	2.7332
month2	1.1674	0.2055	5.6802	0.0000	0.7646	1.5702
month3	2.0868	0.2050	10.181	0.0000	1.6851	2.4885
month4	3.0009	0.1950	15.391	0.0000	2.6187	3.3830
month5	3.1219	0.2494	12.516	0.0000	2.6330	3.6108
month6	2.4022	0.2845	8.4435	0.0000	1.8446	2.9599
month7	1.9796	0.3285	6.0268	0.0000	1.3358	2.6233
month8	2.6224	0.2760	9.5004	0.0000	2.0814	3.1634
month9	3.6537	0.2077	17.595	0.0000	3.2467	4.0607
month10	3.1680	0.1864	16.996	0.0000	2.8027	3.5334
month11	1.3457	0.2530	5.3198	0.0000	0.8499	1.8414
month12	1.1511	0.2299	5.0062	0.0000	0.7004	1.6018
attend_v	0.2263	0.0276	8.1979	0.0000	0.1722	0.2804
Model statistics:						
R-squared		0.9912				
Adj. R-squared		0.9905				
F-statistic		9.098e+04	(Prob F-statistic = 0.0000)			
No. Observations		516				
P-value (F-stat)		0.0000				
Distribution		chi2(37)				
Cov. Estimator		robust				

Problem 2 continued on next page. . .

Python Code

```
endog = data[['attend_v']]
other_controls  = df.select(pl.all().exclude("assaults", "ln_assaults", "wkd_ind", "^atten
exog = data[['attend_m', 'attend_n'] + other_controls]
instruments = data[["attend_v_f", "attend_m_f", "attend_n_f", "attend_v_b", "attend_m_b",
ivolsmod = IV2SLS(dependent=data[["ln_assaults"]], endog=endog, exog=exog, instruments=ins
res_ivols = ivolsmod.fit()
print(res_ivols.summary)
```

(ii) **Answer:** after doing the test we get a p value of 0.000 thus we reject the null hypothesis of that no of the moves have an impact on assault and thus conlude there is significant evidence to conclude that there is an association between the movies and assaults.

Table 19: Linear Equality Hypothesis Test

Statistic	**Value**	**Notes**
Test Statistic	122.5900	-
p-value	0.0000	-
Distribution	chi2(3)	-

Python Code

```
#iii
coeffs = res_ivols.params[['attend_v', 'attend_m', 'attend_n']].values
cov = res_ivols.cov.loc[['attend_v', 'attend_m', 'attend_n'], ['attend_v', 'attend_m', 'at

delta_x = np.array([6, -2, -1])

delta_ln_assaults = np.dot(delta_x, coeffs)

std_error = np.sqrt(np.dot(delta_x, np.dot(cov, delta_x)))

lower = delta_ln_assaults - 1.96 * std_error
upper = delta_ln_assaults + 1.96 * std_error

percent_change = 100 * (np.exp(delta_ln_assaults) - 1)
ci_lower = 100 * (np.exp(lower) - 1)
ci_upper = 100 * (np.exp(upper) - 1)

print(f"Predicted % change in assaults: {percent_change:.2f}%")
print(f"95% CI: [{ci_lower:.2f}%, {ci_upper:.2f}%]")
```

Table 20: Predicted Percentage Change in Assaults

Statistic	**Value**
Predicted % change in assaults	163.90%
95% Confidence Interval (CI)	[97.23%, 253.09%]

(iii) **Python Code**

```python
#ii
print(res_ivols.wald_test(formula="attend_v = attend_m = attend_n = 0"))
```

(e) Based on your analysis, what do you conclude about the effect of violent movies on (short-run) violent behavior?

Answer: yes, but looking at the other results for the other types of movies there seems to an association with the other type of movies and assaults

Problem 3

We examined **Koop and Tobias's data** on wages, education, ability, and so on. We considered the model.

$$\begin{aligned}\ln wage &= \beta_0 + \beta_1 educ + \beta_2 ability + \beta_3 experience \\ &\quad + \beta_4 motherEduc + \beta_5 FatherEduc + \beta_6 broken \\ &\quad + \beta_7 sinlings + \epsilon\end{aligned}$$

(a) We are interested in possible non-linearities in the effect of education on ln Wage. (Koop and Tobias focused on experience. As before, we are not attempting to replicate their results.) A histogram of the education variable shows values from 9 to 20, a spike at 12 years (high school graduation), and a second at 15. Consider aggregating the education variable into a set of dummy variables:

$$\begin{aligned}\text{HS} = 1 &\quad \text{if Educ} \leq 12, \quad 0 \text{ otherwise} \\ \text{Col} = 1 &\quad \text{if } 12 < \text{Educ} \leq 16, \quad 0 \text{ otherwise} \\ \text{Grad} = 1 &\quad \text{if Educ} > 16, \quad 0 \text{ otherwise}\end{aligned}$$

replace Educ in the model with (Col, Grad), making high school (HS) the base category, and recompute the model. Report all results. How do the results change? Based on your results, what is the marginal value of a college degree? What is the marginal impact on ln Wage of a graduate degree?

(b) The aggregation in part (a) actually loses quite a bit of information. Another way to introduce non-linearity in education is through the function itself. Add $educ^2$ to the equation in part (a) and recompute the model. Again, report all results. What changes are suggested? Test the hypothesis that the quadratic term in the equation is not needed—that is, that its coefficient is zero. Based on your results, sketch a profile of log wages as a function of education.

(c) One might suspect that the value of education is enhanced by greater ability. We could examine this effect by introducing an interaction of the two variables in the equation. Add the variable

$$EducAb = Educ \times ability$$

to the base model in part a. Now, what is the marginal value of an additional year of education? The sample mean value of ability is 0.052374. Compute a confidence interval for the marginal impact on ln Wage of an additional year of education for a person of average ability.

(d) Combine the models in (b) and (c). Add both $educ^2$ and $EducAb$ to the base model in the beginning of the question and re-estimate. As before, report all results and describe your findings. If we define low ability as less than the mean and high ability as greater than the mean, the sample averages are −0.798563 for the 7,864 low-ability individuals in the sample and +0.717891 for the 10,055 high-ability individuals in the sample. Using the formulation in part (b), with this new functional form, sketch, describe, and compare the log wage profiles for low- and high-ability individuals.

(e) Suppose that you are now given the following regression model:

$$\begin{aligned}\ln(\text{wage}) = \beta_0 &+ \beta_1 \text{educ} \times \mathbf{1}(\text{educ} < \tau) + \beta_2 \text{educ} \times \mathbf{1}(\text{educ} \geq \tau) + \beta_3 \text{exp} \\ &+ \beta_4 \text{MotherEduc} + \beta_5 \text{FatherEduc} + \beta_6 \text{broken} \\ &+ \beta_7 \text{siblings} + \epsilon\end{aligned}$$

where τ is the threshold parameter, and

$$\mathbf{1}(\text{Educ} < \tau) = \begin{cases} 1 & \text{if Educ} < \tau, \\ 0 & \text{otherwise} \end{cases} \quad \text{and} \quad \mathbf{1}(\text{Educ} \geq \tau) = \begin{cases} 1 & \text{if Educ} \geq \tau, \\ 0 & \text{otherwise.} \end{cases}$$

Problem 4

Using the **California test score date**, estimate the regression below using Nonlinear Least Squares. Report you coefficient estimates and standard errors.

$$TestScore = \beta_0(1 - e^{\beta_1(income - \beta_2)}) + \epsilon$$

Problem 5

Use the **Consumption.xlsx**. We have previously estimated the nonlinear consuption function below using nonlinear least squares in class:

$$C = \alpha + \beta Y^{\gamma} + \epsilon$$

Where C is the real consumption and Y is the real disposable income. Alternatively, we can assume that the error term has a normal distribution and estimate the nonlinear regression above using the maximum likeligood estimation (MLE) approach, In particular, the MLE approach maximizes the log-likelihood function given by:

$$L(\alpha, \beta, \gamma, \sigma^2) = -\frac{n}{2}\log(\sigma^2) - \log(2\pi) - \frac{1}{2\sigma^2}\sum_{i=1}^{n}(C_i - \alpha - \beta Y^{\gamma})^2$$

Where σ^2 is the variance of the error term. Using a statistical programming language of you choice, estimate the regression model using the maximum likelihood estimation approach. Your estimate are expected to be similar to those in Table 7.1 of Green's textbook. Please submit the following:

(a) Your code used to preform the estimation.

(b) The output of your estimation, including the estimated parameters and the error variance.